

2017-07-03

- Initial set up of neural net, based on leNet-5, no signal out, most likely due to using unprocessed data.
- The data is very padded, this has been reduced with the function `cropHeart(inp)`, but I will need to make sure all the files are the same size before they get fed into the CNN.
- I could try normalising the data to get a signal, but will need to get the unpadded data working first.

2017-07-04

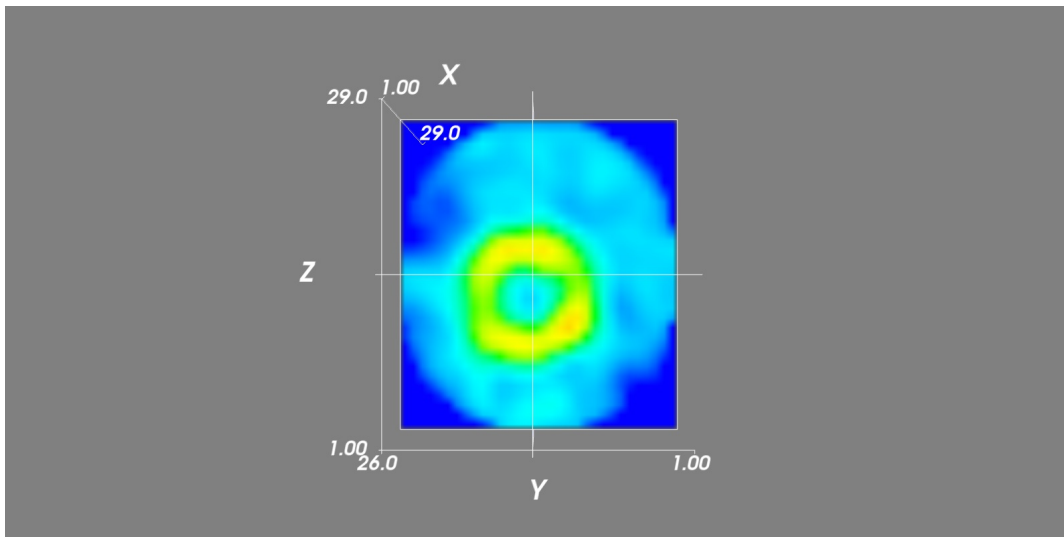
- Wrote a visualisation of the data (`visualisation.py`).
- Still working on repadding the cropped data (It's a bit of a pain).

2017-07-05

- Repadded the cropped data, it is now of size [68,34,34].
- Retrying the CNN with the new data doesn't get a signal. Maybe there isn't enough data to make it work?
- I will fiddle with the hyperparams to see if I can pick something up.
- Maybe normalising the data will help.

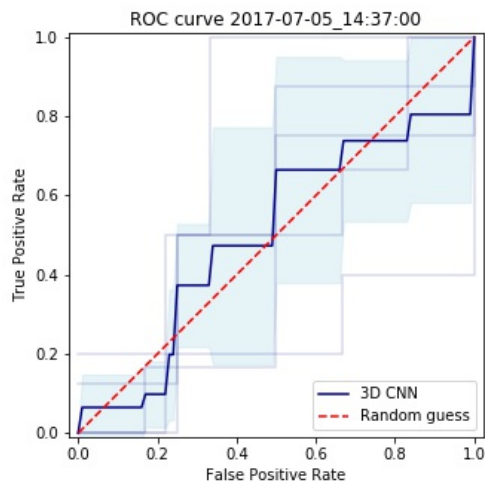
Got some results using 2D slices:

- The 2D slices I used look like:

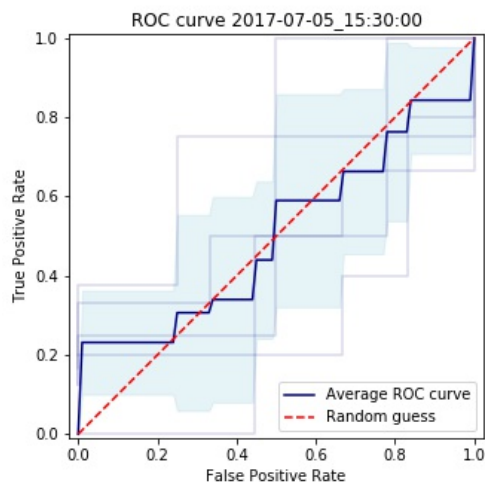


- I have used 2d slices of the data and it works well (halfway through the z-axis). It uses:
 - Slice of rest and slice of stress on z axis. Spatial x and spatial y on x and y axes.
 - LeNet-5 CNN with 3D convolution and subsampling.
 - [2,5,5] filters, pooling 2 with step 2.
 - learning rate of 0.0001, with ADAM optimiser, and batch size of 10.
 - After 50 epochs of 58 images it learns to ~95%.
- I will now apply a k-fold x-validation to it to see if it's not just picking up noise.

- The k=10-fold x-validation shows that the CNN is learning the noise in the data, although this could be due to the small amount of images in each k-fold (only 6!):



- I tried normalising the arrays, with no luck. It stopped overfitting the data, but still hasn't learnt significantly:

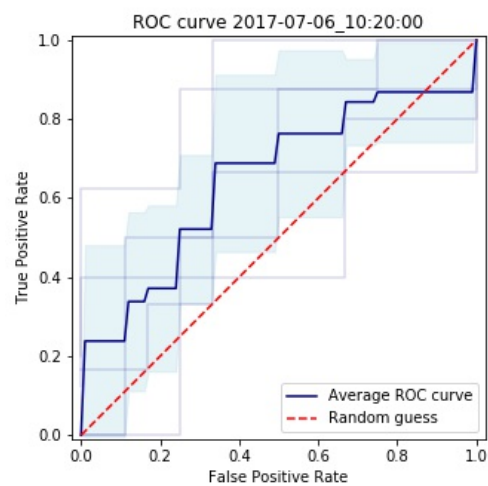
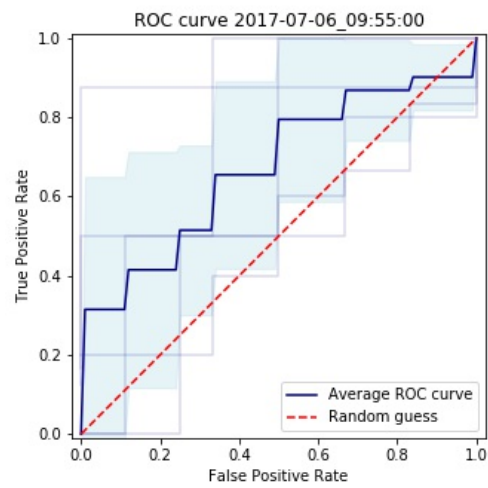
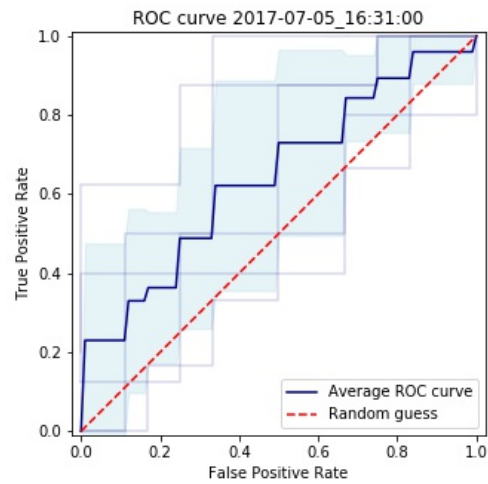


- I think the issue is still the massive amount of blankspace. I should try and scale the arrays so that they are the same size.

Have a signal!

- I have got a signal with the following CNN:
 - Slice of rest and slice of stress on z axis. Spatial x and spatial y on x and y axes.
 - LeNet-5 CNN with 3D convolution and subsampling.
 - [2,10,10] filters, pooling 2 with step 2.
 - learning rate of 0.0001, with ADAM optimiser, and batch size of 10.
 - 5 k-folds.
 - After 50 epochs of 47 images it learns training data to ~95%.
 - Over three repeats:
 - Avg Spec: 0.583, 0.623, 0.663
 - Avg Sens: 0.633, 0.683, 0.700

- ROC curves:



2017-07-06

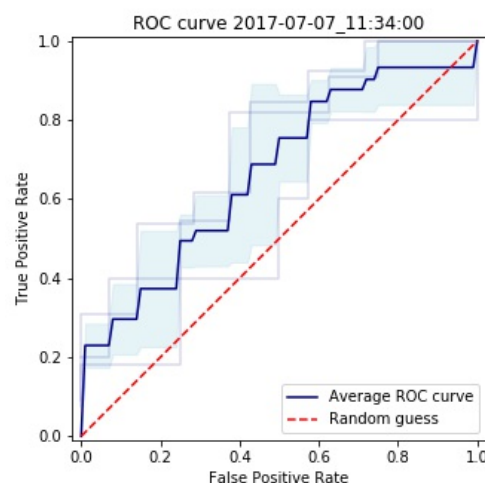
- I redid the 2D slice data with three slices along the x, y, and z axes. It will take ~100 mins to finish learning. It's probably time to use some better hardware.

- Found a function (<https://docs.scipy.org/doc/scipy-0.14.0/reference/generated/scipy.ndimage.interpolation.zoom.html>) which should work well for resizing the images.
- Maybe the reason that the slicing works, and the 3D data doesn't is because the CNN filter only sees one 3D image at a time, and sees both the rest and stress images at the same time in the 2D slice data. I could write a 4D CNN to fix this.
 - mhuen seems to have written a 4D convolution by stacking 3D CNN outputs (<https://github.com/mhuen/TFScripts/blob/master/py/tfScripts.py>). This might work for what I want to do, and stacking can be used for pooling too.
- I wrote a scaling function that eliminates most of the whitespace. After training the CNN did not learn significantly.
- Added a ROC AUC calculator to the outputs.
- I'm going to try artificially expanding the data.

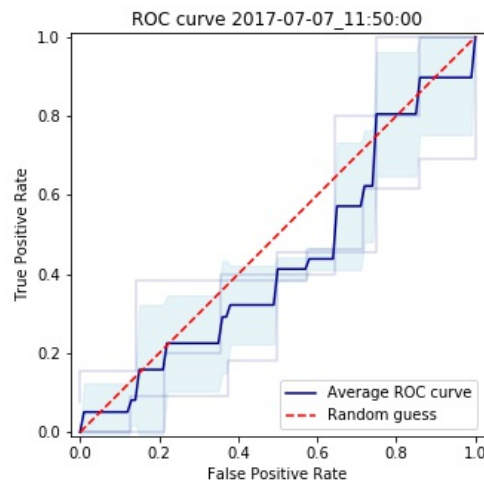
2017-07-07

- Because of the overfitting going on when running the CNN, I increased the L2 regularisers' weight decay from 0.001 to 0.01, and added an extra dropout layer between the two FC dense layers.
- Can't seem to get any results with a spec/sens over 60%, probably due to the way I'm organising the data.
- The CNN appears to train better when using non-scaled data. I can't figure out why. Maybe it's using the image sizes as an aid?
 - Conv filter: [2,15,15]; pool filter: [2,2,2]; 2 FC 1024 neurons, L2 regularisation at weight decay = 0.001, dropout at 0.5 after each FC layer; ADAM optimiser, learning rate = 0.0001, categorical x-entropy loss; batch size = 10 at 38 datum; k = 3 folds.

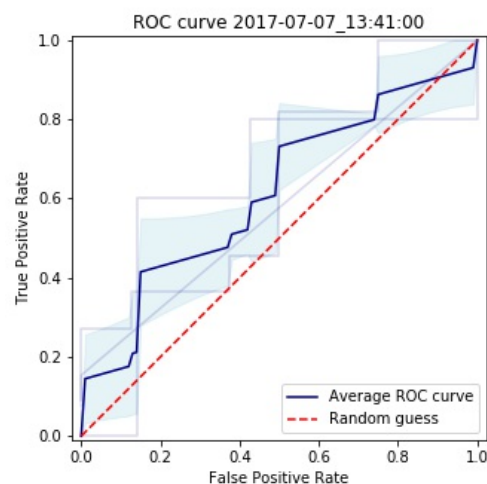
■ non-scaled:



- scaled:



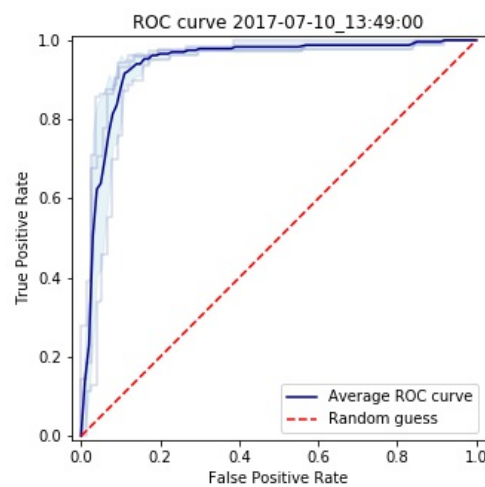
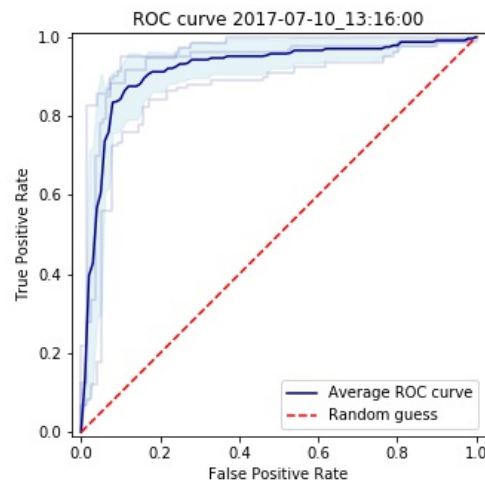
- scaled, not renormalised:



- As shown in the ROC curves, the only data that is causing consistent learning is the non-scaled one. I don't know why.
- Rewrote heart_data.ipynb so that it can resize the input data.

2017-07-10

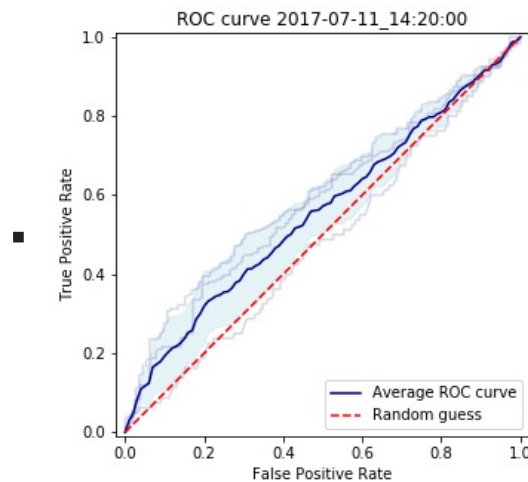
- It might be better to use a Siamese CNN instead of a 4D CNN to compare two 3D images, as training will be faster.
 - I have written CNNs using two-channel, and Siamese architectures, along with the OG 3D convolution architecture. The two-channel and Siamese architectures are described here: <https://arxiv.org/pdf/1504.03641.pdf>.
- The use of a very deep NN architecture would reduce linearity, and may be useful.
- Artificially expanding the data seems to have worked. I am getting after $k = 3$ folds (100 epochs) at 619 datum (two runs):
 - Spec: 0.864, 0.917
 - Sens: 0.888, 0.883
 - ROC AUC: 0.918, 0.940
 - This is with the two-channel architecture. ROC curves:



- Haven't got any significant results from the Siamese CNN, but have only trained it to ~30 epochs. It will probably need more training than the two-channel as there are nearly twice as many weights in the Siamese CNN.
- I should try validating the CNN on ppts that it hasn't seen before (like take 10 ppts from the pool before artificial expansion and use these to validate).

2017-07-11

- I have separated ppts into different k-folds before expansion, so each k-fold has unique ppts in it now, even after artificial expansion. We'll see how it performs now... (This is in the 2channel ipynb)
 - It doesn't work very well. Getting ~50% accuracy.



- More data would be helpful to reduce overfitting, but using all three dimensions of the heart data may be enough to get "good enough" results.
- I have written a 2 channel CNN for the 3D data. It should be ready to try on the supercomputer.

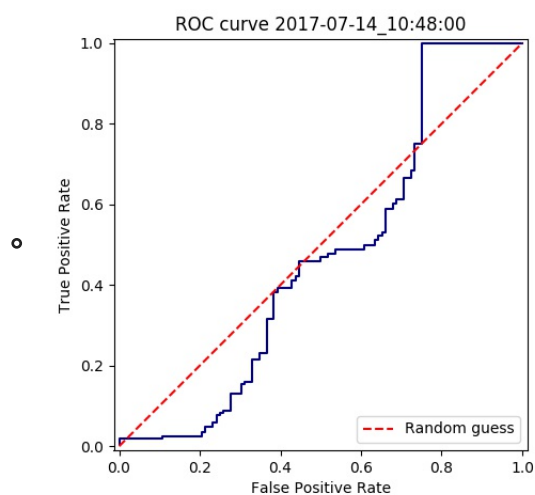
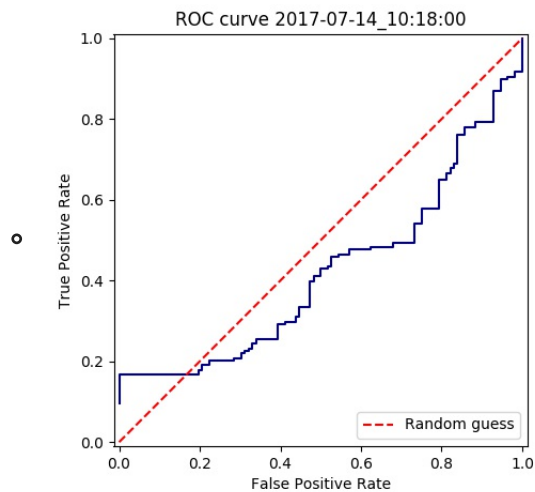
2017-07-12

- Testing the 2dSiameseCNN on the supercomputer:
`$ qsub -q gpu -l nodes=1:ppn=16 -I -X -l walltime=24:00:00`
 It doesn't seem to work. How long is the queue?
- I have found a bug in the 2dsliceCNN that may be causing the lack of learning. The expansion doesn't relabel the expanded data correctly. I have hopefully fixed this.
- Running for 20 epochs @ $k = 5$ folds to see how it does.
 - Again, ~50% accuracy.
- I have increased the number of conv layers to 4.
 - No change.
- Running the 3D CNN on the hub. It looks like it takes ~20 epochs to train to 100% (I should use validation to see if/when it starts overfitting). It also takes ~12s to train an epoch. To contrast it takes my computer ~16mins per epoch, an 80x speedup.
- Added 4(!) new convolution layers to 3D CNN. Since this reduces linearity, we may find something.
 - Getting some odd results. The CNN comes out with the opposite of what I was expecting (low ROC, accuracy).
 - Look at labelling, try on simpler data (MNIST 0s and 1s?), reintroduce k-folding?

2017-07-14

- Added overall average performance metric to 3dCNN-nokfold.

- I think I have found the cause of the low ROC/accuracy. The random state shuffle is set to 1. If I change it, it may get some more believable results.
- Looks like that was what the issue was. The CNN got lucky with the cubes taken out for testing:



- Using fake data to train a CNN. It's found on /data/jim/Heart/sims.
- The CNNs aren't training. For normal/infarction data I have the loss decreasing but the accuracy is static.